

ASR: Past Paper May 2023

Completed on April 24, 2023

Patrick Tourniaire

Problem 1

Subproblem 1. *Why is a pre-emphasis filter of the following form applied to the speech signal?*

$$x[n'] = x[n] + \alpha x[n - 1] \quad (1)$$

Answer

Pre-emphasis is a processing method which increases the amplitude of high frequency bands and decreases the amplitudes of lower bands. The traditional formulation for pre-emphasis is given by Equation 2.

$$x[n'] = x[n] - \alpha x[n - 1] \quad (2)$$

Which measure the transition between signal timesteps, whereas Equation 1 further exaggerates this by not subtracting the previous amplitude. Therefore, the pre-emphasis of this given form is a first order high-pass filter which amplifies the high frequency components of the signal. It achieves this by adding a scaled version of the previous sample to the current sample.

Subproblem 2. *Why is a Hamming window often used rather than a rectangular window? In your answer sketch the shape of each window in the time domain and sketch the effect of each on a sine wave of constant frequency.*

Answer

A Hamming window is often used rather than a rectangular window in signal processing because it helps reduce the effects of spectral leakage and improve the frequency resolution.




Figure 1: Rectangular window effect




Figure 2: Hamming window effect

As we can observe in Figure 1, the rectangular window will potentially introduce a large drop in the signal amplitude which is interpreted as spectral leakage. However, for the Hamming window we do not observe this pattern as highlighted by Figure 2. Since this window filter has a slight taper at the ends as opposed to the drastic drop for the rectangular window, we get a smoother filtering of the sinus wave.

Subproblem 3. *How does a mel-scale filter bank model the human perception of frequency?*

Answer

The mel-scale filter bank model is a model of human auditory perception that aims to simulate the way we hear and process sound by being more discriminative at lower bands and more discriminatory at higher bands. Which mimicks the non-linearity of the human auditory system. In Figure 3 we can observe how this model prioritises lower frequencies.




Figure 3: Mel-scale filter bank model example

Additionally, we observe that the length of the triangles in the lower frequencies are smaller, which will result in more frequencies being canceled and reduced significantly. Whereas the higher frequencies have longer triangles, thus ensuring that a larger range of frequencies are accounted for. Which highlights how model indeed follows the non-linearity of the human auditory system.

Subproblem 4. *What is cepstral analysis and how is it used to separate source and filter characteristics in speech?*

Answer

Cepstral analysis is used as the power spectrum of a signal has some useful properties for separating the source and filter in a signal. We define the source and filter by the following.

- **Source:** refers to the vocal cords and the way they vibrate to produce sound.
- **Filter:** shape and characteristic of the vocal tract that modify the sound produced by the vocal cords.

Cepstral analysis is therefore used to separate the characteristics of source and filter by decomposing the log spectrum of the speech signal into its constituent parts. The cepstrum process looks like the following.

1. Compute the FFT (fast-fourier transform) of the speech signal to obtain its frequency spectrum.
2. Compute the power spectrum of the speech signal by taking the magnitude squared of the FFT.
3. Take the log of the power spectrum to obtain the log spectrum.
4. Take the inverse fourier transform of the log spectrum to obtain the cepstrum.
5. Analyse the cepstrum to separate the source and filter characteristics of the speech signal.

The first few coefficients correspond to the filter characteristics of the speech signal, while the later coefficients correspond to the source characteristics. The plot over the different coefficients can be observed in Figure




Figure 4: MFCC plot with the cepstrum coefficients

Subproblem 5. Consider an HMM with an output distribution given by a (single) Gaussian per state. We would like to estimate the Gaussian mean vector μ_j of state s_j , given a training sequence of acoustic observations $X = x_1, \dots, x_t, \dots, x_T$. In Viterbi training the time-state alignment is assumed to be known. Write an expression for the maximum likelihood estimate of μ_j .

Answer

The expression for the estimated μ_j , given that we have a single Gaussian optimised using MLE is given by.

$$\hat{\mu}_j = \frac{1}{T} \sum_{i=0}^T x_i \quad (3)$$

Subproblem 6. In the case where the time-state alignment is not known, the EM algorithm may be used to estimate the mean vector. The EM algorithm makes use of the state occupation probability of state s_j at time t , $\gamma_t(s_j)$, where s_j is a state of the HMM.

- Express $\gamma_t(s_j)$ as a conditional probability.
- How may $\gamma_t(s_j)$ be written in terms of the forward and backward probabilities?
- Explain how $\gamma_t(s_j)$ may be used to re-estimate μ_j .

Answer

We can express $\gamma_t(s_j)$ as a conditional probability by the following expression.

$$\gamma_t(s_j) = P(q_t = j | X, M) \quad (4)$$

Recalling that $P(X|M) = \alpha_E$ we can express this term by the forward (α) and backward probabilities (β), where M represents the HMM topology.

$$= \frac{1}{\alpha_E} \alpha_j(s_j) \beta_j(s_j) \quad (5)$$

For re-estimating the means of the state Gaussians, we can treat the occupation probabilities as priors representing the likelihood of some RV realisation x belonging to some Gaussian component j . Thus, the estimation for parameter μ becomes.

$$\hat{\mu}_j = \frac{\sum_{t=0}^T \gamma_t(s_j) x_t}{\sum_{t=0}^T \gamma_t(s_j)} \quad (6)$$

Subproblem 7. *Why are HMM computations usually performed in the log domain?*

Answer

HMM computations are performed in the log domain due to the fact that we have to deal with long likelihood chains which are multiplied together to estimate the parameters of the model. However, the problem with this is that this can result in incredibly small numbers which in turn can result in float computation errors when the computer performs the computation. Therefore, we convert to logs as then the multiplication essentially becomes the same as an addition of log likelihoods, as the following property applies.

$$\log(ab) = \log(a) + \log(b) \quad (7)$$

Problem 2

Subproblem 8. *Describe two ways that neural networks trained to perform phone recognition can be used to provide discriminative features for HMM/GMM speech recognition systems (tandem systems).*

Answer

One way neural networks trained for phone recognition can be used to provide discriminative features for HMM/GMM speech recognition systems is through a tandem approach. In this approach, the neural network is used as a front-end feature extractor for the HMM/GMM system. The output of the neural network, which is typically a set of bottleneck features, is concatenated with the traditional MFCC features extracted from the speech signal. The concatenated features are then used as input to the HMM/GMM system for phone recognition.

Another way neural network can be used to provide discriminative features for HMM/GMM speech recognition systems is through a hybrid approach. In this approach, a neural network is used to estimate the posterior probability of the phone labels given the acoustic features. These posteriors are then used as input to the HMM/GMM system as a replacement for the traditional Gaussian mixture model likelihood scores. The HMM/GMM system can then use these posteriors to estimate the most likely phone sequence given the input speech signal. This approach has been shown to improve phone recognition performance, particularly in noisy environments.

Subproblem 9. *Consider a hybrid HMM / neural network system.*

- *How do hybrid systems differ from tandem systems?*
- *In a context-independent hybrid HMM / neural network system, why is the neural network output corresponding to phone q , divided by an estimate of the prior probability $P(q)$?*

A tandem system of a neural network and a HMM/GMM, results in a system where the neural network supplies additional features for the HMM/GMM system to use. Which is done by concatenating the features to the traditional MFCC features. On the other hand, a hybrid approach will simply use the features of the neural network, where it outputs posterior probabilities for the phones given the features.

In a context-independent hybrid HMM/NN system, the neural network is used to estimate the posterior probability of a phone given the acoustic features. The posterior probability of a phone is the probability of the phone given the observed features, and it is proportional to the product of the likelihood of the features given the phone and the prior probability of the phone.

In this context-independent system, the prior probability of each phone is assumed to be equal. However, in reality, some phones occur more frequently than others in the training data. Therefore, dividing the NN output corresponding to phone q by an estimate of the prior probability $P(q)$ helps to normalise the output

and make it less biased towards the more frequent phones. This is known as Bayesian normalisation, and it ensures that the output of the NN is proportional to the likelihood of the features given the phone, and not influenced by the prior probability of the phone.

By normalizing the output in this way, the hybrid system can effectively estimate the posterior probability of a phone given the observed features, which can be used to improve phone recognition accuracy.

Subproblem 10. *Consider a context-dependent hybrid HMM / neural network system.*

- *For practical purposes, a context-dependent HMM/GMM system is required before beginning to train a hybrid system. Why is this so?*
- *Compare and contrast two methods of performing speaker adaptation in a context-dependent hybrid system.*

Answer

A context-dependent HMM/GMM system is required before training a hybrid system because it provides the necessary context-dependent labels that the NN needs to learn the mapping between the acoustic features and the corresponding phone labels. Context-independent HMM/GMM systems assume that each phone has a fixed acoustic representation that is independent of its context. However, in reality, the acoustic realisation of a phone depends on the context in which it appears. For example, the acoustic realisation of the phone /t/ may be different in the word "then" compared to the word "stop". Therefore, context-dependent HMM/GMM systems use different models for each phone depending on the surrounding context, which captures the contextual variations of each phone.

The outputs of a context dependent HMM/GMM system are context-dependent phone labels, which are used to train the NN in the hybrid system. The NN learns to map the context-dependent acoustic features to the corresponding context-dependent phone labels, which improves the accuracy of phone recognition. Therefore, a context-dependent HMM/GMM system is required before training a hybrid system to provide the context-dependent phone labels that the NN needs to learn the mapping between the acoustic features and the corresponding phone labels.

As people have varied pronunciations and different accents, single pre-trained models can struggle to pick up such variances. Therefore, we introduce a technique called speaker adaption. Which for a context-dependent model such as the HMM/GMM can have two approaches for achieving speaker adaption, namely feature-space transforms and model-based adaption. Feature-space transforms involve transforming the input acoustic features of a speaker into a space that is better aligned with the average acoustic characteristics of the training data. One common feature-space transform method is maximum likelihood linear regression (MLLR), which learns a linear transformation matrix to adapt the acoustic model for a specific speaker. MLLR has been shown to be effective for speaker adaptation in context-dependent hybrid systems, especially when a small amount of speaker-specific training data is available.

Model-based adaptation involves adapting the parameters of the acoustic model directly using speaker-specific training data. One common model-based adaptation method is speaker-adaptive training (SAT), which involves retraining the acoustic model using speaker-specific data. SAT has been shown to be effective for speaker adaptation in context-dependent hybrid systems, especially when a large amount of speaker-specific training data is available. A major difference between feature-space transforms and model-based adaptation is the amount of speaker-specific data required. *Feature-space transforms can be effective even with a small amount of speaker-specific training data, while model-based adaptation requires a larger amount of speaker-specific data to be effective. However, model-based adaptation can potentially achieve higher accuracy than feature-space transforms, especially when a large amount of speaker-specific training data is available.*

In summary, both feature-space transforms and model-based adaptation are effective methods for speaker adaptation in context-dependent hybrid systems, but they differ in the amount of speaker-specific data required and the potential accuracy gains. Feature-space transforms can be effective with a small amount of

speaker-specific data, while model-based adaptation requires more data but can potentially achieve higher accuracy.

Subproblem 11. *It might be possible to obtain better acoustic likelihood estimates by combining HMM/GMM and hybrid HMM / neural network systems. How could this be done at the frame level?*

Answer